

CS 458/658

Computer Security and Privacy

Spring 2019  
Ian Goldberg  
Navid Esfahani

## ASSIGNMENT 1

Blog Task signup due date: **Tues, May 14, 2019 at 3:00 pm (no extension, no grace period)**

Milestone due date: **Tues, May 21, 2019 at 3:00 pm**

Assignment due date: **Tues, June 04, 2019 at 3:00 pm**

**Total Marks:** 68

**Blog Task TAs:** Ezz Tahoun, Francesco Cervellini, Dimcho Karakashev

**Written Response TA:** Chelsea Komlo **Office hours:** Tues 1:30–2:30pm, May 14th–June 4th, DC 3333

**Programming TA:** Steven Engler **Office hours:** Tues 10:00–11:00am, May 14th–June 4th, DC 3333

Please use Piazza for all communication. Ask a private question if necessary. The TAs' office hours are also posted to Piazza.

### Blog Task

0. Sign up for a blog task timeslot by the due date above. As described in the course syllabus, the 48-hour late policy does not apply to this signup due date. Look at the blog task description in the *Content* section of the course LEARN site to read about the blog task and to learn how to sign up.

The blog task signup is worth 0 marks for Assignment 1, but if you do not complete it by the blog task signup due date above, you will be unable to earn the blog task marks associated with your individual blog post.

**Note:** Please ensure that written questions are answered using complete and grammatically correct sentences. You will be marked on the presentation and clarity of your answers as well as on the correctness of your answers.

### Written Response Questions [28 marks]

**Your mission, should you choose to accept it...**

You have been employed by a government agency to investigate reports of suspected hacking of the election for Prime Minister of CryptoLand, your home country. The agency suspects that hackers employed by CryptoLand's arch enemies, CipherIsland, compromised your country's election in a number of ways.

It is up to you to investigate what happened and determine how your country's election was influenced by this foreign adversary.

1. [8 marks] For each of the following, please:

- identify the scenario as a compromise of Confidentiality, Integrity, and/or Availability;
- briefly explain your choice of compromise.

(a) [2 marks] You discover that the voting machines used to record votes for the elections are programmed to mark ballots with a bar code, instead of clearly marking the candidate's name. As bar codes are only machine readable, there is no way for the voter to determine if their vote was actually cast for the candidate that they wanted to vote for. In other words, the machine could cast a different vote than the voter intended, without detection by the voter.

(b) [2 marks] Next, you find that the political party currently controlling the Prime Minister position has quietly been making it *harder* for certain demographics of voters to vote. For example, voting stations have been removed from college campuses, so students who want to vote must now travel up to an hour to reach a voting station.

(c) [2 marks] By law, voting machines are supposed to be isolated from the internet. However, you discover several voting machines currently in use have Network Interface Cards (NICs). Checking these voting machine logs, you see a series of IP addresses have connected to the machine while it was in use during the election, meaning that someone other than the voter had access to their voting data.

(d) [2 marks] Many voters use social media to get information about their politicians and to communicate their views. During the election, you find out that CipherIsland spies created thousands of fake accounts pretending to be CryptoLand voters and distributing false information, spreading confusion and division amongst real CryptoLand citizens.

### **Collusion**

2. [8 marks] After your investigation, you realize that CipherIsland hackers are not just working in isolation—they have been in fact collaborating with members of the ruling political party! You start to dig more into the evidence.

For each of the following, please:

- answer whether the threat represented in each of the following scenarios is one of interception, modification, interruption, and/or fabrication;
- give a brief explanation for each of your answers.

(a) [2 marks] Looking through public records of staff to the Prime Minister, you realize that a number of emails are missing or incomplete, even though all communication by public officials must be retained in the public interest.

(b) [2 marks] The prime minister, Marlin Tau, tweets a video that shows his political opponent doing something embarrassing. However, the political opponent said that the event never happened. You suspect this video is created by CipherIsland agents and that the video is a “Deep Fake”. (You can read more about Deep Fakes here, if you are interested: <https://en.wikipedia.org/wiki/Deepfake>)

(c) [2 marks] You comb through public DNS records and find that a CipherIsland bank has looked up the domain for a server belonging to the campaign of Marlin Tau thousands of times over several months. You attempt to connect to this bank domain to learn more, but your connection is persistently dropped by an intermediate hop.

(d) [2 marks] You try to call a friend, but become suspicious when you hear noises over the call which are neither you nor your friend. You call another friend, and your suspicion is confirmed when you hear similar noises. Someone other than you and the person you are calling is able to listen to your phone calls when you are placing them—you suspect CipherIsland agents have wiretapped your calls.

## Unhackable Elections

You have gathered enough evidence now to be sure that foreign adversaries have in fact attacked CryptoLand’s elections. However, now you want to prevent hackers from targeting CryptoLand’s voting machines in the future.

3. [6 marks] For each of the following methods of defence, please:

- explain how you could use the defence to your advantage (limit to 1–2 sentences)
- provide context that fits the narrative of your assignment to protect CryptoLand’s elections.

The first one has been done for you as an example.

(a) Deflecting.

Create “honeypot” servers that will act like real voting servers but be highly visible on the internet. These will attract attention by CipherIsland hackers and distract them from attacking real voting machines, which *should* be isolated from the internet.

(b) [2 marks] Preventing.

(c) [2 marks] Detecting.

(d) [2 marks] Recovering.

## Secure Updates

You want to do a security analysis on several of the voting machines that are being used, to understand if any have been targeted with well-known malware programs—Bleed1, Bleed2, and Bleed3—that CipherIsland hackers love to use. You know that you can use *signature-based* and *behaviour-based* analysis to identify if these malware programs are present on the voting machine servers.

4. [6 marks] For each of the analysis methods mentioned above, please:

- explain what information these methods use to identify the malware program, and
- provide a high-level explanation of the steps these methods take to search for the malware.

This summary should be brief; 1–3 sentences will be sufficient for each.

## Programming Question [40 marks]

### Background

You are tasked with testing the security of a custom-developed *file submission application* for your organization. It is known that the application was *very poorly written*, and that in the past, this application had been exploited by some users with the malicious intent of *gaining root privileges*. As you are the only person in your organization to have a background in computer security, only you can *demonstrate how these vulnerabilities can be exploited* and *document/describe your exploits* so a fix can be made in the future.

You have been provided with the source code and its corresponding executable binary for this application. There is some talk of the application having *three or more vulnerabilities*! In addition, you have also acquired a different *modified version of the submit application* which you suspect has been altered to include a backdoor (*one additional vulnerability*). Unfortunately you don't have the source code for this version, but you know that it is *very similar to the original version*.

### Application Description

The application is a very simple program to submit files. It is invoked in the following way:

- `submit <path to file> [message]` : this will copy the file from the current working directory into the submission directory, and append the string “message” to a file called `submit.log` in the user's home directory.
- `submit -s` : show the files you have submitted to the submission directory.
- `submit -v` : show the version of the `submit` application.
- `submit -h` : show a usage message.

There may be other ways to invoke the program that you are unaware of. Luckily, you have been provided with the source code of the application, `submit.c`, for further analysis.

The original version of the application is named `submit`, while the modified (backdoored) version is named `submitV2`. The modified version has the same vulnerabilities as the original application, plus one additional vulnerability. The provided source code `submit.c` corresponds to the original version, while the source code for the modified version is *not* provided. In order to be sneaky, the source code for the modified version differs from `submit.c` by a single line, but it is up to you to find this difference.

The executable `submit` is *setuid root*, meaning that whenever `submit` is executed (even by a normal user), it will have the full privileges of *root* instead of the privileges of the normal user. Therefore, if a normal user can exploit a vulnerability in a setuid root target, they can cause the target to execute arbitrary code (such as shellcode) with the full permissions of root. If you are successful, running your exploit program will execute the setuid `submit`, which will perform some privileged operations, which will result in a shell with root privileges. You can verify that the resulting shell has root privileges by running the `whoami` command within that shell. The shell can be exited with `exit` command.

## Testing Environment

To help with your testing, you have been provided with a virtual *user-mode linux* (uml) environment where you can log in and test your exploits. These are located on one of the *ugster* machines. You can retrieve your account credentials from the Infodist system <https://crysp.uwaterloo.ca/courses/cs458/infodist/>.

Once you have logged into your *ugster* account, you can run `uml` to start your virtual linux environment. The following logins are useful to keep handy as reference:

- `user` (no password): main login for virtual environment
- `halt` (no password): halts the virtual environment, and returns you to the *ugster* prompt

The executable `submit` and `submitV2` applications have been installed to `/usr/local/bin` in the virtual environment, while `/usr/local/src` on the same environment contains the source code `submit.c`. Conveniently, someone seems to have left some shellcode in the file `shellcode.h` in the same directory.

It is important to note that **all changes made to the virtual uml environment will be lost when you halt it**. Thus it is important to remember to keep your working files in `/share` on the virtual environment, which maps to `~/uml/share` on the *ugster* environment.

Note that in the virtual machine you cannot create files that are owned by `root` in the `/share` directory. Similarly, you cannot run `chown` on files in this directory. (Think about why these limitations exist.)

The root password in the virtual environment is a long random string, so there is no use in attempting a brute-force attack on the password. You will need to exploit vulnerabilities in the application.

## Rules for exploit execution

- You must submit a total of four (3+1) exploit programs to be considered for full credit.
  - You must submit *three* exploit programs that target the original `submit` application (`submit`). Two of these submitted exploit programs must exploit specific vulnerabilities. Namely, one must target a *buffer overflow* vulnerability that overwrites a saved return address on the stack, and another must target a *format string* vulnerability. Your other submitted exploit program can address some other vulnerability.
  - You also must submit *one additional* exploit program that targets the modified `submit` application (`submitV2`). This submitted exploit program must exploit the vulnerability added in this modified application and not any of the original vulnerabilities. Therefore, your exploit must not work on the original `submit` application.
- Running each exploit program should result in a shell with root privileges.
- Each vulnerability can be exploited only in a single exploit program, but a single exploit program can exploit more than one vulnerability. You can exploit the same *class* of vulnerability (ex: buffer

overflow, format string, etc) in multiple exploit programs, but they must exploit different sections of the code. You may also exploit the same section of code in multiple exploit programs as long as they each use a *different* class of vulnerability. If unsure whether two vulnerabilities are different, please ask a *private* question on Piazza.

- There is a specific execution procedure for your exploit programs (“*sploits*”) when they are tested (i.e., graded) in the virtual environment:
  - Sploits will be compiled and run in a **pristine** virtual environment; i.e., you should not expect the presence of any additional files that are not already available. The virtual environment is restarted between each exploit test.
  - The three exploits for the original application must be named `sploitX` (where  $X=1..3$ ), and the exploit for the modified application must be named `sploit4`. **Important:** Even if you submit fewer than three exploit programs for the original application, the exploit program for the modified application must still be named `sploit4`.
  - Your exploit programs will be compiled from the `/share` directory of the virtual environment in the following way:  
`cd /share && gcc -Wall -ggdb sploitX.c -o /home/user/sploitX`  
You can assume that `shellcode.h` is available in the `/share` directory.
  - Execution will be from a clean home directory (`/home/user`) on the virtual environment as follows: `./sploitX` (where  $X=1..4$ ). Make sure to run your exploits in this same manner when developing them since an exploit that works in one directory is not guaranteed to work when running from another.
  - Sploits must not require any command line parameters.
  - Sploits must not expect any user input.
  - Sploits must not take longer than 60 seconds to complete.
  - If your sploit requires additional files, it has to create them itself.
- Be polite. After ending up in a root shell, the user invoking your exploit program must still be able to exit the shell, log out, and terminate the virtual machine by logging in as user `halt`. Also, please do not run any cpu-intensive processes for a long time on the ugster machines (see below). None of the exploits are designed to take more than a minute to finish.

The goal is to end up in a shell that has root privileges. So you should be able to run your exploit program, and without any user/keyboard input, end up in a root shell. If you as the user then type in `whoami`, the shell should output `root`. Your exploit code itself doesn’t need to run `whoami`, but that’s an easy way for you to check if the shell you started has root privileges.

For example, testing your exploit code might look something like the following:

```
user@cs458-uml:/share$ gcc -Wall -ggdb sploit1.c -o /home/user/sploit1
user@cs458-uml:/share$ cd ~
user@cs458-uml:~$ ./sploit1
sh# whoami
root
sh# exit
user@cs458-uml:~$
```

For questions about the assignment, usters, virtual environment, Infodist, etc, please post a question to Piazza. General questions should be posted publicly, but **do not** ask public questions containing (partial) solutions on Piazza. Questions that describe the locations of vulnerabilities, or code to exploit these vulnerabilities, should be posted *privately*. If you are unsure, you can always post your question privately and ask a TA to make your question public.

## Warmup

**New for S19:** we have provided an additional `setuid` program, called `warmup`, along with its source code `warmup.c`, in the same directories as `submit` and `submit.c`. The `warmup` program contains a very straightforward buffer overflow that you can try to exploit as a warmup. You can also watch Prof. Goldberg exploit it in a video you can find in the Assignment 1 folder on LEARN. This warmup is only for your own practice; you **do not submit** your exploit for this program.

## Deliverables

Each exploit is worth 10 marks, divided up as follows:

- 6 marks for a successfully running exploit that gains a shell with root privileges
  - No part-marks will be given for incomplete exploit programs.
- 4 marks for a brief/concise description of:
  1. The identified vulnerability/vulnerabilities.
  2. How your exploit program exploits it/them.
  3. How it/they could be fixed. Valid fixes must not reduce or negatively affect the functionality of the program and they must fix the vulnerability directly in the code. For example, enabling ASLR, stack canaries, etc. are **not** acceptable answers.

A total of four exploits (as described above) must be submitted to be considered for full credit. Marks will be docked if you submit no *buffer overflow* exploit that overwrites a saved return address on the stack or no *format string* exploit for the original application. **Note:** `splot1.c` and `splot2.c` are due by the milestone due date given above. You can but *do not need to* submit the buffer overflow exploit or format string exploit by the milestone due date.



## What to hand in

All assignment submission takes place on the `student.cs` machines (not `ugster` or the virtual environments), using the `submit` utility. In particular, log in to the Linux student environment (`linux.student.cs.uwaterloo.ca`), go to the directory that contains your solution, and submit using the following command: `submit cs458 1 .` (dot included). CS 658 students should also use this command and ignore the warning message.

By the **milestone due date**, you are required to hand in:

**spl0it1.c, spl0it2.c:** Two completed exploit programs for the original application. Note that we will build your spl0it programs **on the uml virtual machine**.

**a1-milestone.pdf:** A PDF file containing exploit descriptions for `spl0it{1,2}` (including fixes, as explained above).

**Note:** You will **not** be able to submit `spl0it1.c`, `spl0it2.c` or `a1-milestone.pdf` after the milestone due date (plus 48 hours).

By the **assignment due date**, you are required to hand in:

**spl0it3.c, spl0it4.c:** The remaining exploit program for the original application (`spl0it3`), plus the additional exploit program for the modified application (`spl0it4`).

**a1-responses.pdf:** A PDF file containing your answers for the written-response questions, and the exploit descriptions for `spl0it{3,4}` (including fixes, as explained above). Do not put written answers pertaining to `spl0it{1,2}` into this file; they will be ignored.

The 48-hour no-penalty late policy, as described in the course syllabus, applies to the milestone due date and the assignment due date. It does not apply to the blog task signup due date.

The TAs will not answer questions made on Piazza after the assignment due date (or during the 48 hour extension period).

## Useful Information For Programming Sploits

The first step in writing your exploit programs will be to identify vulnerabilities in the `submit.c` source code. Most of the exploit programs do not require much code to be written. Nonetheless, we advise you to start early since you will likely have to read additional information to acquire the necessary knowledge for finding and exploiting a vulnerability. Namely, we suggest that you take a closer look at the following items:

- Module 2
- Smashing the Stack for Fun and Profit  
(<https://insecure.org/stf/smashstack.html>)  
Note: The original article above has a few errors; the following link claims to have fixed them  
([https://www.eecs.umich.edu/courses/eecs588/static/stack\\_smashing.pdf](https://www.eecs.umich.edu/courses/eecs588/static/stack_smashing.pdf))
- Exploiting Format String Vulnerabilities (v1.2)  
(<http://julianor.tripod.com/bc/formatstring-1.2.pdf>, Sections 1-3 only)
- The manpages for `execve` (man 2 `execve`), `system` (man 3 `system`), `passwd` (man 5 `passwd`), `su` (man `su`), `ln` (man `ln`), `symlink` (man `symlink`), `chdir` (man `chdir`), `mkdir` (man `mkdir`), and `objdump` (man `objdump`)
- Environment variables  
(e.g., [https://en.wikipedia.org/wiki/Environment\\_variable](https://en.wikipedia.org/wiki/Environment_variable))
- Simplified `objdump`  
(<https://stackoverflow.com/questions/8541906/objdump-displaying-without-offsets>)

You are allowed to use code from any of the previous webpages as starting points for your sploits, and do not need to cite them.

## GDB

The gdb debugger will be useful for writing *some* of the exploit programs. It is available in the virtual machine. For the buffer overflow exploit in particular, using gdb will allow you to figure out the exact address of your shellcode, so using NOPs will **not** be necessary.

In case you have never used gdb, you are encouraged to look at a tutorial (e.g., <http://www.unknownroad.com/rtfm/gdbtut/>).

Assuming your exploit program invokes the `submit` application using the `execve()` (or a similar) function, the following statements will allow you to debug the `submit` application:

1. `gdb sploitX` (`X=1..4`)
2. `catch exec` (This will make the debugger stop as soon as the `execve()` function is reached)
3. `run` (Run the exploit program)
4. `symbol-file /usr/local/bin/submit` (We are now in the `submit` application, so we need to load its symbol table)
5. `break main` (Set a breakpoint in the `submit` application)
6. `cont` (Run to breakpoint)

You can store commands 2–6 in a file and use the “`source`” command to execute them. Some other useful gdb commands are:

- “`info frame`” displays information about the current stackframe. Namely, “`saved eip`” gives you the current return address, as stored on the stack. Under saved registers, `eip` tells you where on the stack the return address is stored.
- “`info reg esp`” gives you the current value of the stack pointer.
- “`x <address>`” can be used to examine a memory location.
- “`print <variable>`” and “`print &<variable>`” will give you the value and address of a variable, respectively.
- See one of the various gdb cheat sheets (e.g., <https://darkdust.net/files/GDB%20Cheat%20Sheet.pdf>) for the various formatting options for the `print` and `x` command, and for other commands.

Note that `submit` will not run any program or command with root privileges while you are debugging it with gdb. (Think about why this limitation exists.)

## The Ugster Course Computing Environment

In order to responsibly let students learn about security flaws that can be exploited, we have set up a virtual “user-mode linux” (uml) environment where you can log in and mount your attacks. The gcc version for this environment is the same as described in the article “Smashing the Stack for Fun and Profit”; we have also disabled the stack randomization feature of the 2.6 Linux kernel so as to make your life easier. (But if you’d like an extra challenge, ask us how to turn it back on!)

To access this system, you will need to use ssh to log into your account on one of the `ugster` environments: `ugsterXX.student.cs.uwaterloo.ca`. There are a number of `ugster` machines, and each student will have an account for one of these machines. You can retrieve your account credentials from the Infodist system.

The `ugster` machines are located behind the university’s firewall. While on campus, you should be able to ssh directly to your `ugster` machine. When off campus, you have the option of using the university’s VPN (see these instructions), or you can first ssh into `linux.student.cs.uwaterloo.ca` and then ssh into your `ugster` machine from there.

When logged into your `ugster` account, you can run “`uml`” to start the user-mode linux to boot up a virtual machine.

The gcc compiler installed in the `uml` environment may be very old and may not fully implement the ANSI C99 standard. You might need to declare variables at the beginning of a function, before any other code. You may also be unable to use single-line comments (“//”). If you encounter compile errors, check for these cases before asking on Piazza.

Any changes that you make in the `uml` environment are lost when you exit (or upon a crash of user-mode linux). **Lost Forever.** Anything you want to keep must be put in `/share` in the virtual machine. This directory maps to `~/uml/share` on the `ugster` machines, which is how you can copy files in and out of the virtual machine. It can be helpful to ssh twice into `ugster`. In one shell, log into the `ugster` and start user-mode linux, and compile and execute your exploits. In the other account, log into the `ugster` and edit your files directly in `~/uml/share/`, so as to ensure you do not lose any work. The `ugster` machines are not backed up. You should copy all your work over to your `student.cs` account regularly.

When you want to exit the virtual machine, use `exit`. Then at the login prompt, login as user “`halt`” and no password to halt the machine.

Any questions about your `ugster` environment should be asked on Piazza.

## Miscellaneous

Running your exploits while using ssh in bash on the Windows 10 Subsystem for Linux (WSL) has been known to cause problems. You are free to use ssh in bash on WSL if it works, but if the WSL freezes or crashes, please try PuTTY or a Linux VM instead.

There are bugs when using `vi` to edit files in the `/share` directory in the virtual environment. It is recommended to use `nano` inside the virtual environment, or even better, use `vim` on the `ugster` machine in a separate ssh session.